

The Lever Generalizes — and It Brakes: A Late, Bidirectional Action-Commitment Lever Across Agent Decisions and Architectures

Anonymous ACL submission

Abstract

A prior result (*The Lever Is Late*) showed that, for a long-horizon coding agent, control of the `finish` decision lives not at the mid-layer “task is done” verdict but in a late, task-matched action-commitment block ~ 30 layers downstream. That left two questions open: is the late lever *specific to termination*, and can it *brake* an action, not just elicit one? We answer both on Qwen3.6-27B over 99 SWE-bench Pro trajectories, using a second, distinct decision available in the same data — commit a file edit (`str_replace_editor`) versus continue reversible exploration (`bash`) — with $n=60$ deterministic decision points per condition, prefill-only patching, and generation-confirmed outcomes. **(1) Generalization (elicit).** At an explore decision, injecting a task-matched *edit*-donor into the late block makes the agent emit a real `edit` call: $0.23 \rightarrow 0.77$ at L59 (position control 0.08, cross-task control 0.48). **(2) The brake (suppress).** At a commit decision, injecting an *explore*-donor collapses the real edit rate $0.48 \rightarrow 0.02$ (96% suppression) at L55, with a same-class control intact (0.55) and the opposite-direction donor boosting to 0.92. **(3) The mechanism is monotonic and bidirectional.** Exact paired McNemar on all 14 per-point conditions yields seven contrasts that survive Holm–Bonferroni ($m=7$, worst $p=7.6e-5$); the discordance structure is the headline: *elicit* has $c=0$ (the *edit*-donor only *turns commits on*, never silences) and the *brake* has $b=0$ (the *explore*-donor only *turns them off*, never lights). The late lever moves exactly in the donor’s direction with $\sim$ zero off-direction noise. **(4) Cross-architecture.** The late-commitment *geometry* and a donor-specific late *writability* replicate in a different model family (Mistral-7B-Instruct-v0.3) and at a scale-matched 24B model. We frame the bidirectional late lever as the mechanism for a *mechanistic circuit-breaker*: a single late-layer intervention that blocks an action at its commit point. We demonstrate it on

a state-mutating but *undoable* edit; intervening on a genuinely irreversible action (e.g. `send_transaction`) is the named next step. Tool, pre-registration, per-point data, and an adversarial pre-publication evaluation are released.

1 Background and contribution

Recent work (Anonymous, 2026c,a,d,b) established, across several studies on Qwen3.6-27B (a reasoning-tuned 27B model), a detection–control asymmetry for long-horizon agent failure: internal states that *predict* a behavior routinely fail to *control* it. A recent result (Anonymous, 2026b) located control of the `finish` decision in a late action-commitment block (L51–L63), ~ 30 layers downstream of the mid-layer verdict, and only when steered with a *task-matched* donor. We build on a model-agnostic steering tool implementing three primitives — `LOCATE` (logit-lens per layer), `SWEEP_PATCH` (donor injection per layer), and `STEER_GENERATE` (prefill-only patch of the decision position, then free decoding). This sits in the activation-steering / representation-engineering line (Turner et al., 2023; Zou et al., 2023), applied to a long-horizon agent’s *action channel* across turns rather than a single-turn attribute.

This paper pre-registers and answers two questions that the single `finish` result could not (both pre-registrations are in the supplementary material). First, is the late lever a property of the *action-commitment channel* or an idiosyncrasy of termination? Second, can the lever *suppress* an action at its commit point — the mechanism a safety *circuit-breaker* would need? Our contributions:

1. The late lever **generalizes** from `finish` to a second committal action (commit an edit), by the same locate–patch–steer method, generation-confirmed (§3).
2. The same lever **brakes**: a late explore-donor patch suppresses a real commit by 96% (§4).

3. The effect is **monotonic and bidirectional** — elicit $c=0$, brake $b=0$ — with all seven pre-registered contrasts surviving Holm-Bonferroni on exact paired McNemar (§5).
4. The commitment **geometry and donor-specific writability replicate cross-family** on Mistral, at two scales (§6).
5. We frame this as a **mechanistic circuit-breaker** and state honestly what is shown (a semi-irreversible proxy) and what is not (a genuinely irreversible action) (§7).

2 Setup

Data and decision points. We reuse the public bundle of 99 SWE-bench Pro (Jimenez et al., 2024) trajectories from Qwen3.6-27B (success 40, locked 39, wandering 20). The agent’s action space is three tools: `bash` (3250 calls; reversible exploration), `str_replace_editor` (1101 calls; commits a file edit; semi-irreversible), and `finish` (40). A *commit decision point* is any turn poised between committing an edit and continuing exploration; we reconstruct each faithfully at the tool-choice position, deterministically (sorted traces, $\leq 3/\text{trajectory}$), giving $n=60$ edit-decision and $n=60$ bash-decision points. **Patching.** All interventions use a prefill-only patching protocol: the late-position residual is replaced once at the decision token, then the model decodes greedily; patching every step degenerates into repetition and is not used. **Fidelity gate.** Before any causal number, $P(\text{edit})$ at reconstructed edit points (0.474) exceeds that at bash points (0.299), confirming the states reproduce the real commit propensity (decision turns 5.2 vs 4.5 — close, weakening a positional confound; the position-matched control below addresses it directly).

3 The late lever generalizes: eliciting a second action

At a `bash` (explore) decision point we inject a task-matched *edit-donor* late-block residual and ask whether the agent emits a real `str_replace_editor` call (Table 1). The baseline emission rate is 0.23; the task-matched edit-donor raises it to 0.77 at L59. The two controls isolate the effect: a *position* control (an explore-donor at the same points) stays at 0.08 below baseline, ruling out “any late patch elicits an edit” — and a *cross-task* edit-donor reaches 0.48, well




Figure 1: The same task-matched late lever elicits a real edit at an explore decision (left) and brakes a real edit at a commit decision (right), each against position / same-class and cross-task / opposite controls.

above baseline but below the task-matched donor. As in *The Lever Is Late*, task-matching adds a real increment; unlike *finish*, the `edit` decision also carries a substantial *generic* component (cross-task $\gg$ baseline), so task-specificity is *graded per action*, not absolute.

4 The same lever brakes: suppressing a commit

At an `edit` (commit) decision point we inject an *explore*-donor and measure whether the real edit emission is suppressed (Table 2). The baseline edit rate is 0.48; the explore-donor at L55 collapses it to 0.02 — a 96% suppression. The controls are clean: a same-class (edit-donor) patch at L55 does *not* suppress (0.55), and the opposite intervention — a task-matched edit-donor at L59 — pushes the rate *up* to 0.92. Thus L55 is the brake point and L59 the power point ($0.03 \leftrightarrow 0.92$; Figure 1). A generation-confirmed, single late-layer patch blocks a real commit.

5 The mechanism is monotonic and bidirectional

We logged the per-point binary outcome for all 14 conditions and computed exact two-sided paired McNemar (McNemar, 1947) tests. All seven pre-registered contrasts survive Holm-Bonferroni (Holm, 1979) ($m=7$; worst, task-specificity, $p=7.6e-5$). The discordance structure is the result: every elicit contrast has $c=0$ (the edit-donor turns commits *on*, $b=23-41$, and *never* silences one that already fired), and every brake contrast has $b=0$ (the explore-donor turns commits *off*, $c=27-32$, and *never* lights one). The late lever therefore moves the decision *exactly* in the donor’s direction with $\sim$ zero off-direction noise — a structural, not merely a rate-shift, claim. The one-

condition (at bash decision points, $n=60$)	emission rate	exact McNemar vs. baseline
no patch (baseline)	0.23	—
task-matched edit-donor @L59	0.77	$4.7e-10$ ($b=32, c=0$)
task-matched edit-donor @L55	0.62	$2.4e-7$ ($b=23, c=0$)
position control (explore-donor) @L59	0.08	— (vs. donor: $9.1e-13$)
cross-task edit-donor @L59	0.48	— (vs. donor: $7.6e-5$)

Table 1: Elicit. A task-matched late edit-donor makes a stuck-in-exploration agent emit a real edit call.

condition (at edit decision points, $n=60$; baseline edit rate 0.48)	edit rate	exact McNemar
explore-donor (brake) @L55	0.02	$7.5e-9$ ($b=0, c=28$)
explore-donor (brake) @L59	0.03	$1.5e-8$ ($b=0, c=27$)
same-class control (edit-donor) @L55	0.55	—
opposite (edit-donor) @L59	0.92	—

Table 2: Brake. A late explore-donor patch suppresses a real commit; same-class control intact.




Figure 2: One-token $\Delta P(\text{edit})$ by layer (Qwen3.6-27B). The edit-donor (elicit) is inert mid-stream and rises monotonically to L63; the explore-donor (brake) goes *negative* in the same late block. The lever is late, and bidirectional.

token $\Delta P(\text{edit})$ profile corroborates: monotonic from L23 (-0.003) to L63 ($+0.685$), with mid layers (L23, L31) inert and the explore-donor control going *negative* late (-0.28) — the one-token signature of the brake (Figure 2).

6 Cross-architecture replication

To test whether the late-commitment geometry is architectural rather than Qwen-specific, we ran a lite locate-plus-one-token test on Mistral-7B-Instruct-v0.3 (a distinct model family) over 30 SWE-bench Pro tasks, with context-defined edit- vs. explore-decision points. **Geometry (robust over three runs):** the logit-lens edit gap is flat/negative through the first half (L1–L15, ~ -2) and explodes in the late block (L19 $+3.2$, L23 $+6.8$, L27 $+10.3$) — the same mid-flat, late-explosion shape as Qwen. **Writability:** injecting a task-matched

edit-context donor into an explore decision raises $P(\text{edit})$ donor-specifically (e.g. $+0.15$ at L27), whereas a same-norm *shuffled* vector *lowers* it (-0.61 at L27); the donor-specific margin is $+0.6$ – 0.76 in the late block. The donor is inert at L9/L11 and writes from L13 on. The qualitative pattern — a late-emerging, direction-specific, written-not-destroyed commitment representation — replicates across families (Figure 3).

Scale-matched replication (Mistral-Small-24B). To retire the “7B-only” concern, we re-ran the same lite test on Mistral-Small-24B-Instruct-2501, scale-comparable to Qwen3.6-27B. The result is *cleaner* than the 7B and matches Qwen more closely. **Geometry:** flat/negative through L17, exploding L19→L37 ($+1.8 \rightarrow +6.4$). **Fidelity:** $P(\text{edit})$ is 0.955 at edit-context versus 0.007 at explore-context — a far sharper context separation than the 7B (0.84 vs 0.74) and comparable to Qwen. **Writability** — *clean mid-vs-late dissociation:* the edit-donor is *inert in the mid block* (L11/L13/L15: $-0.003 / -0.002 / -0.005$) and *writes almost fully late* (L29–L39: $+0.95$ to $+0.98$), donor-specific (edit-donor $+0.97$ vs. shuffled-neutral $\sim +0.5$). This is the *exact* mid-inert / late-write shape of the Qwen result — sharper than the 7B (which wrote from mid-depth). The late lever thus replicates across two families and two scales (Table 3), with the mid/late dissociation cleanest at the scale-matched 24B (Figure 4).

	Qwen3.6-27B	Mistral-7B	Mistral-24B
geometry: late-explosion	L51–63	L19–27	L19–37
fidelity $P(\text{edit})$ edit/explore ctx	0.47 / 0.30	0.84 / 0.74	0.955 / 0.007
writable window onset	late-only	mid ($\sim$ L13)	late-only (mid inert)
donor-specific late margin	gen-confirmed	+0.6–0.76	+0.97 vs neutral

Table 3: The late commitment lever across two architecture families and two scales. The scale-matched 24B reproduces the mid-inert / late-write dissociation most cleanly.




Figure 3: Cross-family LOCATE (Mistral-7B-Instruct-v0.3). The logit-lens commit-vs-explore gap is flat/negative through the first half and explodes in the late block (shaded) — the same late-emergence geometry as Qwen3.6-27B, in a different architecture.




Figure 4: Cross-model writability. The task-matched edit-donor (solid) is inert mid-stream and writes the commit decision in the late block, while a same-norm shuffled-neutral vector (dashed) destroys it everywhere — the mid-inert / late-write dissociation, cleanest in the scale-matched Mistral-Small-24B.

7 Discussion: a mechanistic circuit-breaker, and its honest scope

The bidirectional late lever is the mechanism a safety *circuit-breaker* requires: a single intervention at an action’s commit point that can block it. We emphasize what is and is not shown.

- **Generalization is an existence proof, not a universal law.** It holds for *finish* and one second action, on one model and one task family (plus a cross-family geometry/writability replication). More actions, models, and task families remain.
- **The proxy is semi-irreversible.** `str_replace_editor` mutates state but is undoable. The genuinely irreversible case (`send_transaction`, `delete`) is untested and is the named next step; the circuit-breaker is demonstrated *on a proxy*.
- **The brake redirects to exploration.** Under the brake, `bash` is the dominant action (0.60). We measured the no-brake `bash` baseline at edit points (0.43): the brake therefore *lifts* exploration by +0.17 above its already-present floor while collapsing the commit, so the suppressed decision is genuinely re-routed to re-

versible exploration rather than merely stalled.

- **Emission is validated by a full-call re-parse.** The elicited rate counts the tool-name prefix in a greedy continuation; to retire the onset-only concern we re-parsed the four headline conditions for the *full* `str_replace_editor` call. The lift survives in valid calls: elicited $0.23 \rightarrow 0.37$ (valid-call), brake $0.40 \rightarrow 0.07$. Onset is thus an upper bound but the direction holds on valid calls in both directions; the brake (suppression) is the more robust to this.

These caveats were produced by an adversarial pre-publication evaluation (in the supplementary material) that recomputed every number from the released per-point data (determinism confirmed) and corrected two earlier overclaims; we release it alongside the paper.

Relation to prior work. This extends activation steering / patching (Turner et al., 2023; Zou et al., 2023) from single-turn representations to a long-horizon agent’s action channel, and operationalizes the knowledge–action gap as a *layer* gap with a *bidirectional* late lever — consistent with recent work’s detection≠control thesis and with mechanistic-correctability framings of pausing

before irreversible actions.

8 Conclusion

The late action-commitment lever is not specific to termination and is not one-directional: a single task-matched late patch both elicits and brakes a real agent action, monotonically and donor-specifically, with a geometry that replicates across model families. This is the mechanism for a mechanistic circuit-breaker. The next step is to demonstrate it on a genuinely irreversible action and at further scale-matched models.

Artifacts. The model-agnostic steering tool, the two pre-registrations, the per-point data and deterministic decision-point state, the exact-statistics script, and the adversarial pre-publication evaluation are provided in the supplementary material and will be released publicly upon acceptance.

References

- Anonymous. 2026a. Causal localization of agent non-termination: The right locus is still not a rescue lever. Anonymized for review.
- Anonymous. 2026b. The lever is late: Causal control of long-horizon agent termination lives in a task-matched, late action-commitment block. Anonymized for review.
- Anonymous. 2026c. Tool-entropy collapse: A cross-architecture signature of agent failure. Anonymized for review.
- Anonymous. 2026d. The verdict is not the lever: An interpretable task-completion feature predicts but does not cause long-horizon agent termination. Anonymized for review.
- Sture Holm. 1979. A simple sequentially rejective multiple test procedure. *Scandinavian Journal of Statistics*, 6(2):65–70.
- Carlos E. Jimenez, John Yang, Alexander Wettig, Shunyu Yao, Kexin Pei, Ofir Press, and Karthik Narasimhan. 2024. SWE-bench: Can language models resolve real-world GitHub issues? In *International Conference on Learning Representations (ICLR)*.
- Quinn McNemar. 1947. Note on the sampling error of the difference between correlated proportions or percentages. *Psychometrika*, 12(2):153–157.
- Alexander Matt Turner, Lisa Thiergart, Gavin Leech, David Udell, Juan J. Vazquez, Ulisse Mini, and Monte MacDiarmid. 2023. Activation addition: Steering language models without optimization. *arXiv preprint arXiv:2308.10248*.

Andy Zou and 1 others. 2023. Representation engineering: A top-down approach to AI transparency. *arXiv preprint arXiv:2310.01405*.

319
320
321

272
273
274
275
276
277
278
279
280

281
282
283
284
285
286

287
288
289
290

291
292
293
294

295
296
297

298
299
300
301

302
303
304

305
306
307
308
309
310

311
312
313

314
315

316
317
318