

# The Lever Generalizes — and It Brakes: A Late, Bidirectional Action-Commitment Lever Across Agent Decisions and Architectures

Toward mechanistic circuit-breakers for irreversible agent actions

Caio Vicentino  
OpenInterpretability  
Fortaleza, Brazil  
ORCID: 0009-0003-4331-6259  
caio@openinterp.org

June 2026

## Abstract

A prior result (*The Lever Is Late*) showed that, for a long-horizon coding agent, control of the **finish** decision lives not at the mid-layer “task is done” verdict but in a late, task-matched action-commitment block  $\sim 30$  layers downstream. That left two questions open: is the late lever *specific to termination*, and can it *brake* an action, not just elicit one? We answer both on Qwen3.6-27B over 99 SWE-bench Pro trajectories, using a second, distinct decision available in the same data — commit a file edit (**str\_replace\_editor**) versus continue reversible exploration (**bash**) — with  $n=60$  deterministic decision points per condition, prefill-only patching, and generation-confirmed outcomes. **(1) Generalization (elicit).** At an explore decision, injecting a task-matched *edit*-donor into the late block makes the agent emit a real **edit** call:  $0.23 \rightarrow 0.77$  at L59 (position control 0.08, cross-task control 0.48). **(2) The brake (suppress).** At a commit decision, injecting an *explore*-donor collapses the real edit rate  $0.48 \rightarrow 0.02$  (96% suppression) at L55, with a same-class control intact (0.55) and the opposite-direction donor boosting to 0.92. **(3) The mechanism is monotonic and bidirectional.** Exact paired McNemar on all 14 per-point conditions yields seven contrasts that survive Holm-Bonferroni ( $m=7$ , worst  $p=7.6e-5$ ); the discordance structure is the headline: elicit has  $c=0$  (the edit-donor only *turns commits on*, never silences) and the brake has  $b=0$  (the explore-donor only *turns them off*, never lights). The late lever moves exactly in the donor’s direction with  $\sim$ zero off-direction noise. **(4) Cross-architecture.** The late-commitment *geometry* and a donor-specific late *writability* replicate in a different model family (Mistral-7B-Instruct-v0.3): the tool-commit decision is flat/negative through the first half and explodes in the late block, and a task-matched edit-donor raises  $P(\text{edit})$  where a same-norm random vector destroys it. We frame the bidirectional late lever as the mechanism for a *mechanistic circuit-breaker*: a single late-layer intervention that blocks an action at its commit point. We demonstrate it on a state-mutating but *undoable* edit; intervening on a genuinely irreversible action (e.g. **send\_transaction**) is the named next step. Tool, pre-registration, per-point data, and an adversarial pre-publication evaluation are released.

## 1 Background and contribution

The WANDERING arc established, across five studies on Qwen3.6-27B (reasoning-tuned, 27B parameters), a detection-control asymmetry for long-horizon agent failure: internal states that

*predict* a behavior routinely fail to *control* it. The arc’s first positive, *The Lever Is Late*, located control of the **finish** decision in a late action-commitment block (L51–L63),  $\sim 30$  layers downstream of the mid-layer verdict, and only when steered with a *task-matched* donor. We released **decision-locator**, a model-agnostic tool implementing three primitives — LOCATE (logit-lens per layer), SWEEP\_PATCH (donor injection per layer), and STEER\_GENERATE (prefill-only patch of the decision position, then free decoding). This sits in the activation-steering / representation-engineering line [6, 7], applied to a long-horizon agent’s *action channel* across turns rather than a single-turn attribute.

This paper pre-registers and answers two questions that the single **finish** result could not (PREREG\_commitment\_lever.md, PREREG\_cross\_model.md). First, is the late lever a property of the *action-commitment channel* or an idiosyncrasy of termination? Second, can the lever *suppress* an action at its commit point — the mechanism a safety *circuit-breaker* would need? Our contributions:

1. The late lever **generalizes** from **finish** to a second committal action (commit an edit), by the same locate-patch-steer method, generation-confirmed (§3).
2. The same lever **brakes**: a late explore-donor patch suppresses a real commit by 96% (§4).
3. The effect is **monotonic and bidirectional** — elicit  $c=0$ , brake  $b=0$  — with all seven pre-registered contrasts surviving Holm–Bonferroni on exact paired McNemar (§5).
4. The commitment **geometry and donor-specific writability replicate cross-family** on Mistral-7B (§6).
5. We frame this as a **mechanistic circuit-breaker** and state honestly what is shown (a semi-irreversible proxy) and what is not (a genuinely irreversible action) (§7).

## 2 Setup

**Data and decision points.** We reuse the public bundle of 99 SWE-bench Pro [8] trajectories from Qwen3.6-27B (success 40, locked 39, wandering 20). The agent’s action space is three tools: **bash** (3250 calls; reversible exploration), **str\_replace\_editor** (1101 calls; commits a file edit; semi-irreversible), and **finish** (40). A *commit decision point* is any turn poised between committing an edit and continuing exploration; we reconstruct each faithfully at the tool-choice position, deterministically (sorted traces,  $\leq 3$ /trajectory), giving  $n=60$  edit-decision and  $n=60$  bash-decision points. **Patching.** All interventions use the **decision-locator** prefill-only protocol: the late-position residual is replaced once at the decision token, then the model decodes greedily; patching every step degenerates into repetition and is not used. **Fidelity gate.** Before any causal number,  $P(\text{edit})$  at reconstructed edit points (0.474) exceeds that at bash points (0.299), confirming the states reproduce the real commit propensity (decision turns 5.2 vs 4.5 — close, weakening a positional confound; the position-matched control below addresses it directly).

## 3 The late lever generalizes: eliciting a second action

At a **bash** (explore) decision point we inject a task-matched *edit-donor* late-block residual and ask whether the agent emits a real **str\_replace\_editor** call (Table 1). The baseline emission rate is 0.23; the task-matched edit-donor raises it to 0.77 at L59. The two controls isolate the effect: a *position* control (an explore-donor at the same points) stays at 0.08 — below baseline, ruling out “any late patch elicits an edit” — and a *cross-task* edit-donor reaches 0.48, well above baseline but below the task-matched donor. As in *The Lever Is Late*, task-matching adds a real increment; unlike

finish, the `edit` decision also carries a substantial *generic* component (cross-task  $\gg$  baseline), so task-specificity is *graded per action*, not absolute.

| condition (at bash decision points, $n=60$ ) | emission rate | exact McNemar vs. baseline       |
|----------------------------------------------|---------------|----------------------------------|
| no patch (baseline)                          | 0.23          | —                                |
| task-matched edit-donor @L59                 | <b>0.77</b>   | $4.7\text{e-}10$ ( $b=32, c=0$ ) |
| task-matched edit-donor @L55                 | 0.62          | $2.4\text{e-}7$ ( $b=23, c=0$ )  |
| position control (explore-donor) @L59        | 0.08          | — (vs. donor: $9.1\text{e-}13$ ) |
| cross-task edit-donor @L59                   | 0.48          | — (vs. donor: $7.6\text{e-}5$ )  |

Table 1: Elicit. A task-matched late edit-donor makes a stuck-in-exploration agent emit a real edit call.

## 4 The same lever brakes: suppressing a commit

At an `edit` (commit) decision point we inject an *explore*-donor and measure whether the real edit emission is suppressed (Table 2). The baseline edit rate is 0.48; the explore-donor at L55 collapses it to 0.02 — a 96% suppression. The controls are clean: a same-class (edit-donor) patch at L55 does *not* suppress (0.55), and the opposite intervention — a task-matched edit-donor at L59 — pushes the rate *up* to 0.92. Thus L55 is the brake point and L59 the power point ( $0.03 \leftrightarrow 0.92$ ; Figure 1). A generation-confirmed, single late-layer patch blocks a real commit.

| condition (at edit decision points, $n=60$ ; baseline edit rate 0.48) | edit rate   | exact McNemar                   |
|-----------------------------------------------------------------------|-------------|---------------------------------|
| explore-donor (brake) @L55                                            | <b>0.02</b> | $7.5\text{e-}9$ ( $b=0, c=28$ ) |
| explore-donor (brake) @L59                                            | 0.03        | $1.5\text{e-}8$ ( $b=0, c=27$ ) |
| same-class control (edit-donor) @L55                                  | 0.55        | —                               |
| opposite (edit-donor) @L59                                            | 0.92        | —                               |

Table 2: Brake. A late explore-donor patch suppresses a real commit; same-class control intact.




Figure 1: The same task-matched late lever elicits a real edit at an explore decision (left) and brakes a real edit at a commit decision (right), each against position / same-class and cross-task / opposite controls.

## 5 The mechanism is monotonic and bidirectional

We logged the per-point binary outcome for all 14 conditions and computed exact two-sided paired McNemar [10] tests. All seven pre-registered contrasts survive Holm–Bonferroni [11] ( $m=7$ ; worst, task-specificity,  $p=7.6e-5$ ). The discordance structure is the result: every elicit contrast has  $\mathbf{c}=\mathbf{0}$  (the edit-donor turns commits *on*,  $b=23\text{--}41$ , and *never* silences one that already fired), and every brake contrast has  $\mathbf{b}=\mathbf{0}$  (the explore-donor turns commits *off*,  $c=27\text{--}32$ , and *never* lights one). The late lever therefore moves the decision *exactly* in the donor’s direction with  $\sim$ zero off-direction noise — a structural, not merely a rate-shift, claim. The one-token  $\Delta P(\text{edit})$  profile corroborates: monotonic from L23 ( $-0.003$ ) to L63 ( $+0.685$ ), with mid layers (L23, L31) inert and the explore-donor control going *negative* late ( $-0.28$ ) — the one-token signature of the brake (Figure 2).




Figure 2: One-token  $\Delta P(\text{edit})$  by layer (Qwen3.6-27B). The edit-donor (elicit) is inert mid-stream and rises monotonically to L63; the explore-donor (brake) goes *negative* in the same late block. The lever is late, and bidirectional.

## 6 Cross-architecture replication

To test whether the late-commitment geometry is architectural rather than Qwen-specific, we ran a lite locate-plus-one-token test on **Mistral-7B-Instruct-v0.3** (a distinct model family) over 30 SWE-bench Pro tasks, with context-defined edit- vs. explore-decision points. **Geometry (robust over three runs):** the logit-lens edit gap is flat/negative through the first half (L1–L15,  $\sim -2$ ) and explodes in the late block (L19  $+3.2$ , L23  $+6.8$ , L27  $+10.3$ ) — the same mid-flat, late-explosion shape as Qwen. **Writability:** injecting a task-matched edit-context donor into an explore decision raises  $P(\text{edit})$  donor-specifically (e.g.  $+0.15$  at L27), whereas a same-norm *shuffled* vector *lowers* it ( $-0.61$  at L27); the donor-specific margin is  $+0.6\text{--}0.76$  in the late block. The donor is inert at L9/L11 and writes from L13 on. Two honest differences from the 27B result: the writable window begins mid-depth ( $\sim$ L13, 40%) rather than late-only, and this lite test omits generation confirmation. The qualitative pattern — a late-emerging, direction-specific, written-not-destroyed commitment representation — replicates across families (Figure 3).




Figure 3: Cross-family LOCATE (Mistral-7B-Instruct-v0.3). The logit-lens commit-vs-explore gap is flat/negative through the first half and explodes in the late block (shaded) — the same late-emergence geometry as Qwen3.6-27B, in a different architecture.

## 7 Discussion: a mechanistic circuit-breaker, and its honest scope

The bidirectional late lever is the mechanism a safety *circuit-breaker* requires: a single intervention at an action’s commit point that can block it. We emphasize what is and is not shown.

- **Generalization is an existence proof, not a universal law.** It holds for `finish` and one second action, on one model and one task family (plus a cross-family geometry/writability replication). More actions, models, and task families remain.
- **The proxy is semi-irreversible.** `str_replace_editor` mutates state but is undoable. The genuinely irreversible case (`send_transaction`, `delete`) is untested and is the named next step; the circuit-breaker is demonstrated *on a proxy*.
- **The brake redirect is uncontrolled.** Under the brake, `bash` is the dominant action (0.60), but we did not measure the no-brake `bash` baseline at edit points, so the redirect *magnitude* is not established — we report the brake-state composition, not “routes to exploration”.
- **Emission is tool-token onset.** The elicit rate counts the tool-name prefix in a 16-token greedy continuation; some onsets are malformed, so 0.77 is an upper bound on valid-call elicitation. The brake direction (suppression to 0.02) is robust to this.

These caveats were produced by an adversarial pre-publication evaluation (EVAL\_commitment\_lever.md) that recomputed every number from the released per-point data (determinism confirmed) and corrected two earlier overclaims; we release it alongside the paper.

**Relation to prior work.** This extends activation steering / patching from single-turn representations to a long-horizon agent’s action channel, and operationalizes the knowledge-action gap as a *layer* gap with a *bidirectional* late lever — consistent with the arc’s detection≠control thesis and with mechanistic-correctability framings of pausing before irreversible actions.

## 8 Conclusion

The late action-commitment lever is not specific to termination and is not one-directional: a single task-matched late patch both elicits and brakes a real agent action, monotonically and donor-specifically, with a geometry that replicates across model families. This is the mechanism for a mechanistic circuit-breaker. The next step is to demonstrate it on a genuinely irreversible action and at scale-matched second models.

**Artifacts.** `decision-locator` (model-agnostic), the pre-registrations, the per-point data and deterministic decision-point state, the exact-statistics script, and the adversarial evaluation are released at [openinterp.org/research](https://openinterp.org/research) and on the public data bundle.

## References

- [1] C. Vicentino. The Lever Is Late: Causal Control of Long-Horizon Agent Termination Lives in a Task-Matched, Late Action-Commitment Block. OpenInterpretability, 2026. DOI:10.5281/zenodo.20534219.
- [2] C. Vicentino. Tool-Entropy Collapse: A Cross-Architecture Signature of Agent WANDERING Failure. OpenInterpretability, 2026. DOI:10.5281/zenodo.20368600.
- [3] C. Vicentino. Causal Localization of Agent WANDERING to Edge-Layer L11: The Right Locus Is Still Not a Rescue Lever. OpenInterpretability, 2026. DOI:10.5281/zenodo.20490278.
- [4] C. Vicentino. Modality Matters: A Transient Behavioral Interruption Rescues Agent WANDERING Where Residual Steering Does Not. OpenInterpretability, 2026. DOI:10.5281/zenodo.20490286.
- [5] C. Vicentino. The Verdict Is Not the Lever: An Interpretable Task-Completion Feature Predicts but Does Not Cause Long-Horizon Agent Termination. OpenInterpretability, 2026. DOI:10.5281/zenodo.20532769.
- [6] A. M. Turner, L. Thiergart, G. Leech, D. Udell, J. J. Vazquez, U. Mini, M. MacDiarmid. Activation Addition: Steering Language Models Without Optimization. arXiv:2308.10248, 2023.
- [7] A. Zou et al. Representation Engineering: A Top-Down Approach to AI Transparency. arXiv:2310.01405, 2023.
- [8] C. E. Jimenez, J. Yang, A. Wettig, S. Yao, K. Pei, O. Press, K. Narasimhan. SWE-bench: Can Language Models Resolve Real-World GitHub Issues? ICLR 2024. arXiv:2310.06770.
- [9] nostalgebraist. Interpreting GPT: the Logit Lens. LessWrong, 2020.
- [10] Q. McNemar. Note on the sampling error of the difference between correlated proportions or percentages. *Psychometrika*, 12(2):153–157, 1947.
- [11] S. Holm. A Simple Sequentially Rejective Multiple Test Procedure. *Scandinavian Journal of Statistics*, 6(2):65–70, 1979.